

TUTORIAL PROJECT: MRA MATRIX-MATRIX MULTIPLICATION

Instructions : Complete the tutorial using either C.

To use C, edit only the file C/main.cpp.

See C/mainTemplate.cpp for skeleton code or use to restart if needed.

1 Initialization and Finalization

In C: Set up MPI by using the lines below to Initialize and Finalize MPI in the main(...) method.

```
1 MPI_Init(&argc, &argv);
2 MPI_Finalize();
```

2 Implement Blocked Matrix Matrix Multiply with MPI

In C: We already implemented this method in Homework 1. The only modification that needs to be made to convert blocked_matmat() method to be compatible with MPI.

```
1 void blocked_matmat (int n, double* A, double* B, double* C, int n_iter)
```

3 Implement Fox's and Cannon's Algorithm's

In C: Use your code from Homework 2. Review your code to check for correctness and make edit if needed. Use the function signatures below.

```
1 void fox_matmat (double* A, double* B, double* C, int n, int sq_num_procs,
                  int rank_row, int rank_col)
2 void cannon_matmat (double* A, double* B, double* C, int n, int sq_num_procs,
                      int rank_row, int rank_col)
```

4 Implement RMA Methods

In C: Implement a RMA version of Blocked Matrix-Matrix Multiplication and either a RMA version of Fox's Algorithm or a RMA version Cannon's Algorithms. Use the function signatures below.

```
1 void rma_blocked (int n, double* A, double* B, double* C, int n_iter)
2 void rma_fox (double* A, double* B, double* C, int n, int sq_num_procs,
               int rank_row, int rank_col)
3 void rma_cannon (double* A, double* B, double* C, int n, int sq_num_procs,
                  int rank_row, int rank_col)
```

See Hints for some RMA tips.

5 Compile code

Compile your code with:

```
1 sh compile.sh
```

The script should compile and run all the code. with the following lines as tests.

```

1 //Smaller Tests
2 srun -n 1 ./test 2048
3 srun -n 4 ./test 2048
4 srun -n 16 ./test 2048
5 //Larger Tests
6 srun -n 4 ./test 4096
7 srun -n 16 ./test 4096

```

6 Check for Correctness

Make sure the code runs and gives an output, This is the output from my code for a small test.

```

1 Running: srun -n 1 ./test 2048
2 srun: job 134909 queued and waiting for resources
3 srun: job 134909 has been allocated resources
4 To request GPUs, add --gpus-per-node X or --gpus X, where X is the desired
   number of GPUs.
5 Job 134909 running on easley022
6 MPI Blocked      2.418642 sec
7 MPI Fox          45.085509 sec
8 MPI Cannon        36.801065 sec
9 RMA Blocked       2.446727 sec
10 RMA Cannon        2.772050 sec
11
12 Done

```

Explain the Runtime.

7 Some Hints:

Helpful Slides

Lecture Slides from U-Illinois.

Another Slide Deck from UI.

See "Class 20: One-Sided MPI" lecture slide on Canvas

Useful API's

APIs I used in my RMA implementations

```

1 int MPI_Win_create (void *base, MPI_Aint size, int disp_unit, MPI_Info
  info, MPI_Comm comm, MPI_Win *win)

```

Lock and Unlock

Use Lock and Unlock to use One-Sided Asynchronous Communication. Allows for Root process (R) to make a call to another process (P) without that process (P) sending back info, behaves like shared- memory models.

```

1 MPI_Win_lock(int locktype, int rank, int assert, MPI_Win win)
2 MPI_Win_unlock(int rank, MPI_Win win)

```

Get Put and Free

You only need to use one of the models, Get moves the data early. Put moves the data late. I used get in my implementation.

Get Moves the data: Target → Root.

Put Moves opposite: Root → Target.

```
1 MPI_Get(void *origin_addr, int origin_count, MPI_Datatype origin_dtype,
        int target_rank, MPI_Aint target_disp, int target_count, MPI_Datatype
        target_dtype, MPI_Win win)
2 MPI_Put(void *origin_addr, int origin_count, MPI_Datatype origin_dtype,
        int target_rank, MPI_Aint target_disp, int target_count, MPI_Datatype
        target_dtype, MPI_Win win)
```

Use Free to unallocate the window object, allocation is done in 'create' so don't worry about allocation.

```
1 MPI_Win_free (MPI_Win win)
```

Other Tips

Only consider a square number of processes [2,4,8..], to make the MPI and RMA method better for testing. This allows for perfect partitioning for NxN matrices where N is a perfect square.

Use helper methods since the methods for Matrix Matrix Multiplication have similar implementations. I used `local_accum` and `coords_to_rank` as helpers.

Use `#include <algorithm>` and `#include <vector>` header to use vectors and copy data. Helpful with data flow between buffers. See links below.

```
#include algorithm  
#include vector
```